

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from sklearn import preprocessing
```

```
In [2]: df = pd.read_csv("C:\\Users\\Student\\Desktop\\ajinkya mote- 24\\diabetes.csv")
```

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Pregnancies      768 non-null   int64
1   Glucose          768 non-null   int64
2   BloodPressure    768 non-null   int64
3   SkinThickness    768 non-null   int64
4   Insulin          768 non-null   int64
5   BMI              768 non-null   float64
6   Pedigree         768 non-null   float64
7   Age              768 non-null   int64
8   Outcome          768 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

```
In [5]: df.describe
```

Out[5]:

	<bound method NDFrame.describe of	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	\
0	6	148	72	35	0	33.6		
1	1	85	66	29	0	26.6		
2	8	183	64	0	0	23.3		
3	1	89	66	23	94	28.1		
4	0	137	40	35	168	43.1		
..	...	...	...	...	...	...		
763	10	101	76	48	180	32.9		
764	2	122	70	27	0	36.8		
765	5	121	72	23	112	26.2		
766	1	126	60	0	0	30.1		
767	1	93	70	31	0	30.4		

	Pedigree	Age	Outcome
0	0.627	50	1
1	0.351	31	0
2	0.672	32	1
3	0.167	21	0
4	2.288	33	1
..	...	...	...
763	0.171	63	0
764	0.340	27	0
765	0.245	30	0
766	0.349	47	1
767	0.315	23	0

[768 rows x 9 columns]>

```
In [9]: df.head()
```

Out[9]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [10]: df.corr().style.background_gradient(cmap='BuGn')
```

Out[10]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	Pedigree	Age	Outcome
Pregnancies	1.000000	0.129459	0.141282	-0.081672	-0.073535	0.017683	-0.033523	0.544341	0.221898
Glucose	0.129459	1.000000	0.152590	0.057328	0.331357	0.221071	0.137337	0.263514	0.466581
BloodPressure	0.141282	0.152590	1.000000	0.207371	0.088933	0.281805	0.041265	0.239528	0.065068
SkinThickness	-0.081672	0.057328	0.207371	1.000000	0.436783	0.392573	0.183928	-0.113970	0.074752
Insulin	-0.073535	0.331357	0.088933	0.436783	1.000000	0.197859	0.185071	-0.042163	0.130548
BMI	0.017683	0.221071	0.281805	0.392573	0.197859	1.000000	0.140647	0.036242	0.292695
Pedigree	-0.033523	0.137337	0.041265	0.183928	0.185071	0.140647	1.000000	0.033561	0.173844
Age	0.544341	0.263514	0.239528	-0.113970	-0.042163	0.036242	0.033561	1.000000	0.238356
Outcome	0.221898	0.466581	0.065068	0.074752	0.130548	0.292695	0.173844	0.238356	1.000000

In [11]: df.drop(['BloodPressure', 'SkinThickness'], axis=1, inplace=True)

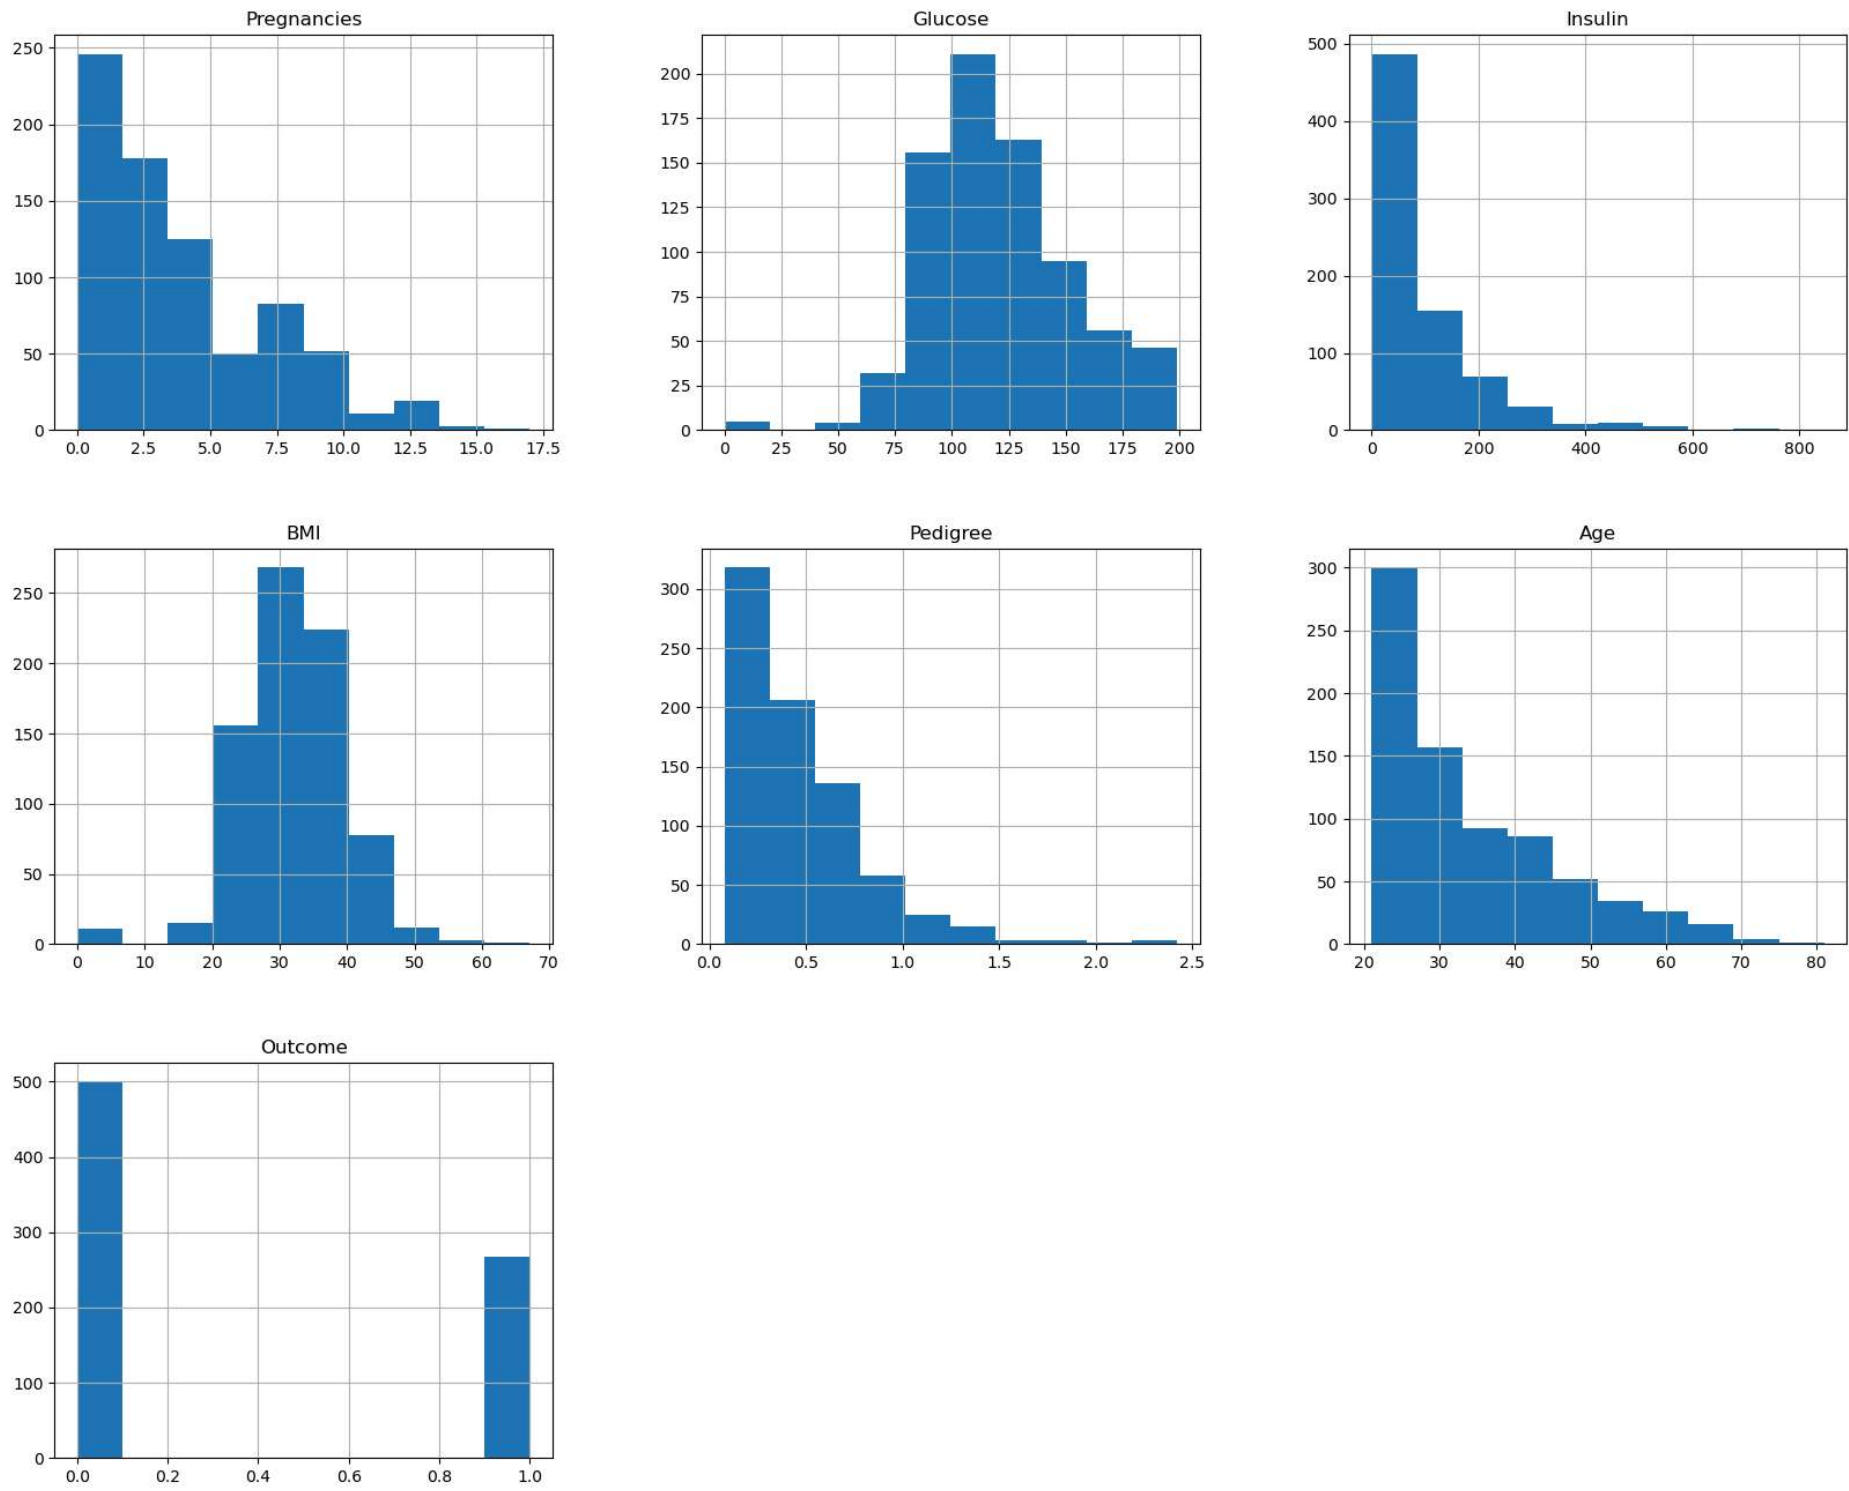
In [12]: df.isna().sum()

Out[12]: Pregnancies 0  
Glucose 0  
Insulin 0  
BMI 0  
Pedigree 0  
Age 0  
Outcome 0  
dtype: int64

In [13]: df.describe

Out[13]: <bound method NDFrame.describe of  
0 6 148 0 33.6 0.627 50 1  
1 1 85 0 26.6 0.351 31 0  
2 8 183 0 23.3 0.672 32 1  
3 1 89 94 28.1 0.167 21 0  
4 0 137 168 43.1 2.288 33 1  
..  
763 10 101 180 32.9 0.171 63 0  
764 2 122 0 36.8 0.340 27 0  
765 5 121 112 26.2 0.245 30 0  
766 1 126 0 30.1 0.349 47 1  
767 1 93 0 30.4 0.315 23 0  
  
[768 rows x 7 columns]>

In [14]: hist = df.hist(figsize=(20,16))



```
In [15]: X=df.iloc[:, :df.shape[1]-1] #Independent Variables
y=df.iloc[:, -1] #Dependent Variable
X.shape, y.shape
```

Out[15]: ((768, 6), (768,))

```
In [16]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=8)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [23]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

def knn(X_train, X_test, y_train, y_test, neighbors, power):
    model = KNeighborsClassifier(n_neighbors=neighbors, p=power)

    # Fit the model using the training set and make predictions on the test set
    y_pred = model.fit(X_train, y_train).predict(X_test)

    # Print accuracy
    accuracy = accuracy_score(y_test, y_pred)
    print(f"Accuracy for K-Nearest Neighbors model \t: {accuracy:.4f}")

    # Compute and print confusion matrix
    cm = confusion_matrix(y_test, y_pred)
    print(f'''Confusion matrix :
    | Positive Prediction\t| Negative Prediction
    +-----+-----+
    Positive Class | True Positive (TP) {cm[0, 0]}\t| False Negative (FN) {cm[0, 1]}
    +-----+-----+
    Negative Class | False Positive (FP) {cm[1, 0]}\t| True Negative (TN) {cm[1, 1]}\n''')

    # Compute and print classification report
    cr = classification_report(y_test, y_pred)
    print('Classification report : \n', cr)
```

```
In [24]: param_grid = {
    'n_neighbors': range(1, 51),
    'p': range(1, 4)
}
grid = GridSearchCV(estimator=KNeighborsClassifier(), param_grid=param_grid, cv=5)
grid.fit(X_train, y_train)
grid.best_estimator_, grid.best_params_, grid.best_score_
```

Out[24]: (KNeighborsClassifier(n\_neighbors=27),  
{'n\_neighbors': 27, 'p': 2},  
0.7719845395175262)

In [27]: knn(X\_train, X\_test, y\_train, y\_test, grid.best\_params\_['n\_neighbors'], grid.best\_params\_['p'])

Accuracy for K-Nearest Neighbors model : 0.7987  
Confusion matrix :  
| Positive Prediction | Negative Prediction  
-----+-----  
Positive Class | True Positive (TP) 91 | False Negative (FN) 11  
-----+-----  
Negative Class | False Positive (FP) 20 | True Negative (TN) 32

Classification report :  
precision recall f1-score support

0	0.82	0.89	0.85	102
1	0.74	0.62	0.67	52
accuracy			0.80	154
macro avg	0.78	0.75	0.76	154
weighted avg	0.79	0.80	0.79	154

In [ ]: